

# **SE31520: Enhancing the CS-Alumni Application**

Due on Monday, December 10, 2012

Alexander Brown

## Contents

<b>Introduction</b>	<b>3</b>
<b>Server Architecture</b>	<b>4</b>
<b>Client Architecture</b>	<b>6</b>
<b>Test Strategy</b>	<b>8</b>
<b>Evaluation</b>	<b>9</b>

## Introduction

The CS-Alumni (CSA) Application prototype is a Ruby on Rails application which runs as a web service. It supports browser-based interaction as well as some RESTful interaction. The purpose of this project is to increase the capabilities of the CSA Application (to prevent confusion I will now refer to it as the CSA server) so it can support a desktop-based client (which I will refer to as the CSA client).

In recent years REST interoperability has increased in popularity and prevalence and, for this sort of project, REST is the best technique for the job. Other methods of interoperability like sockets or RPC would be too heavyweight for purpose, and would just complicate the interaction between client and server.

REST allows a lot of code reuse in the server, especially with Ruby on Rails which has nice ideas about routing and responding formats. As Rails uses all forms of REST requests under the covers, not just the ones supported by browsers, it makes the change from a web-based application to a RESTful service very easy (in fact a lot of it has already been done).

The only thing I really anticipate having to do to the server is expose a little more information for more complex operations such as logging-in or checking access. As these operations are not handled in a RESTful way to make using the web-based version sane (i.e. using sessions to persist log-in credentials).

The only thing the client will need to do is expose REST requests in a nice format to non-technical users. Of course this will also expose the server to other potential clients other than the one discussed here. But again this is not a bad thing, so long as security is maintained. An example of this is a technical user using cURL to issue REST requests via a shell.

## Server Architecture

Having worked in a Level 3 Service team, my attitude toward changing existing code is not to do so unless it is completely necessary and at that, to modify the code as little as possible.

The first change I made was to make a simple addition to the UserController which would echo the username of the currently logged-in user. This allowed me to *a*) confirm the user has been logged-in; *b*) check if the user is an administrator (currently this is defined by the username being equal to “admin”).

This change involved adding a single method (and associated GET route) which only rendered JSON (i.e. it would only accept content-type JSON).

The only other change I made to add functionality to include the details of a user’s image in the JSON file from the GET request to `/user/:id.json`.

Figure 1 shows the overall class diagram for the CSA Server.

As is apparent from figure 1, the architecture of the server hasn’t changed at all. The only additions are those mentioned above (the addition of a single GET route, and some additional information in the User’s JSON file). Other than this I felt no real need to change the architecture of the server, nor the way in which it performs the current functionality.

I did have to fix some small pieces of code in the way emails were produced and sent to make that code actually work. This was part of fixing the tests on the CSA server.

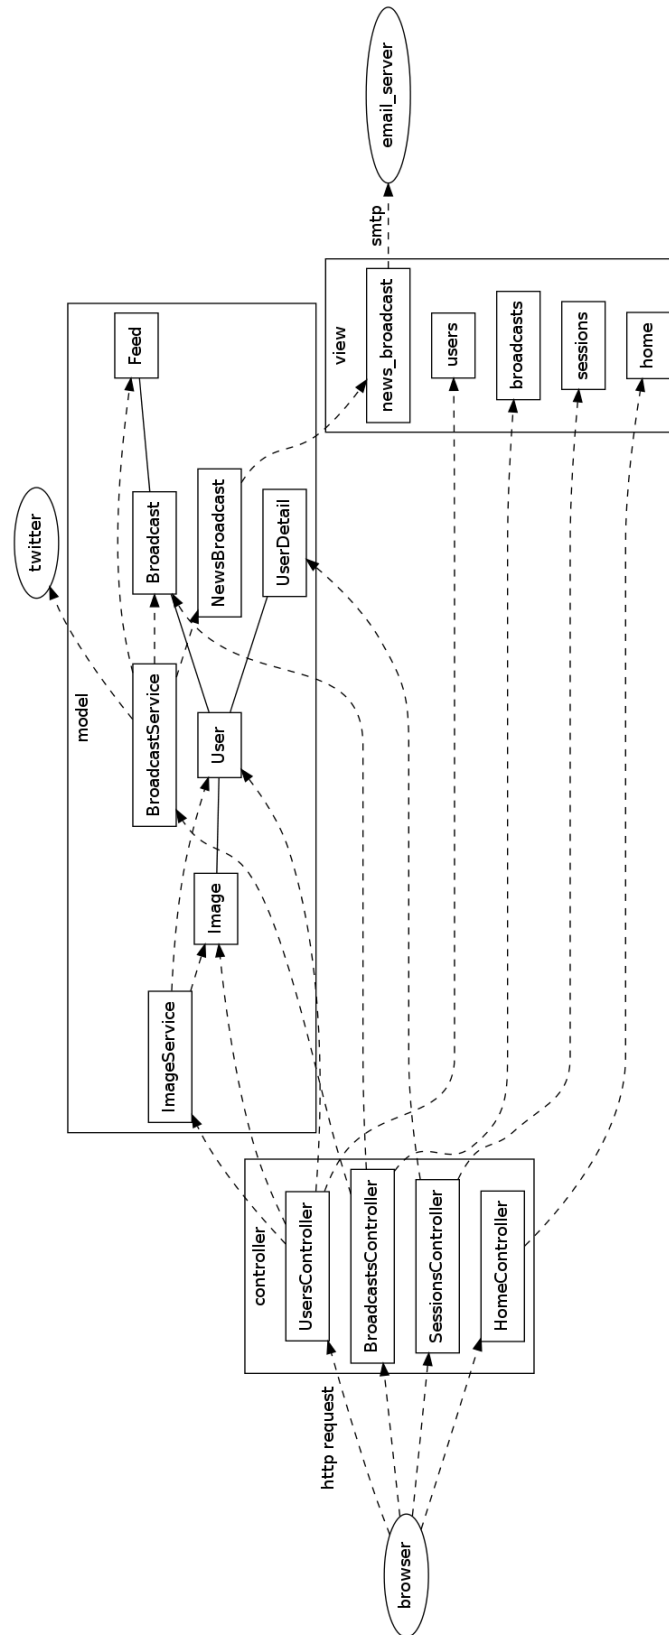


Figure 1: Class diagram for the CSA server

## Client Architecture

As the client doesn't need to do a lot I decided to keep it fairly simple. I still tried to follow the Model, View, Controller (MVC) design pattern as best as I could. Separating out a lot of the REST interaction into a sole class.

This also means it would be relatively easy to swap out this method of interaction without too much effort. The views take up the majority of the class diagram as there are different ways of presenting the same information. However even these have been kept as simple as possible.

Figure 2 shows the UML Class diagram for the client.

From this I then started to consider the choice of language. Ruby was the obvious choice to keep a bit of compatibility between client and server. I also found that Rails has a gem which provides client-based REST access to a Rails application; ActiveSupport. This saved me a lot of time and effort during implementation, not having to work out the exact, correct, REST requests and the form the JSON files passed between the two take.

Ruby also has some nice graphics libraries, FXRuby is one which is cross-platform, with a lot of good documentation, based on the FOX toolkit for C++.

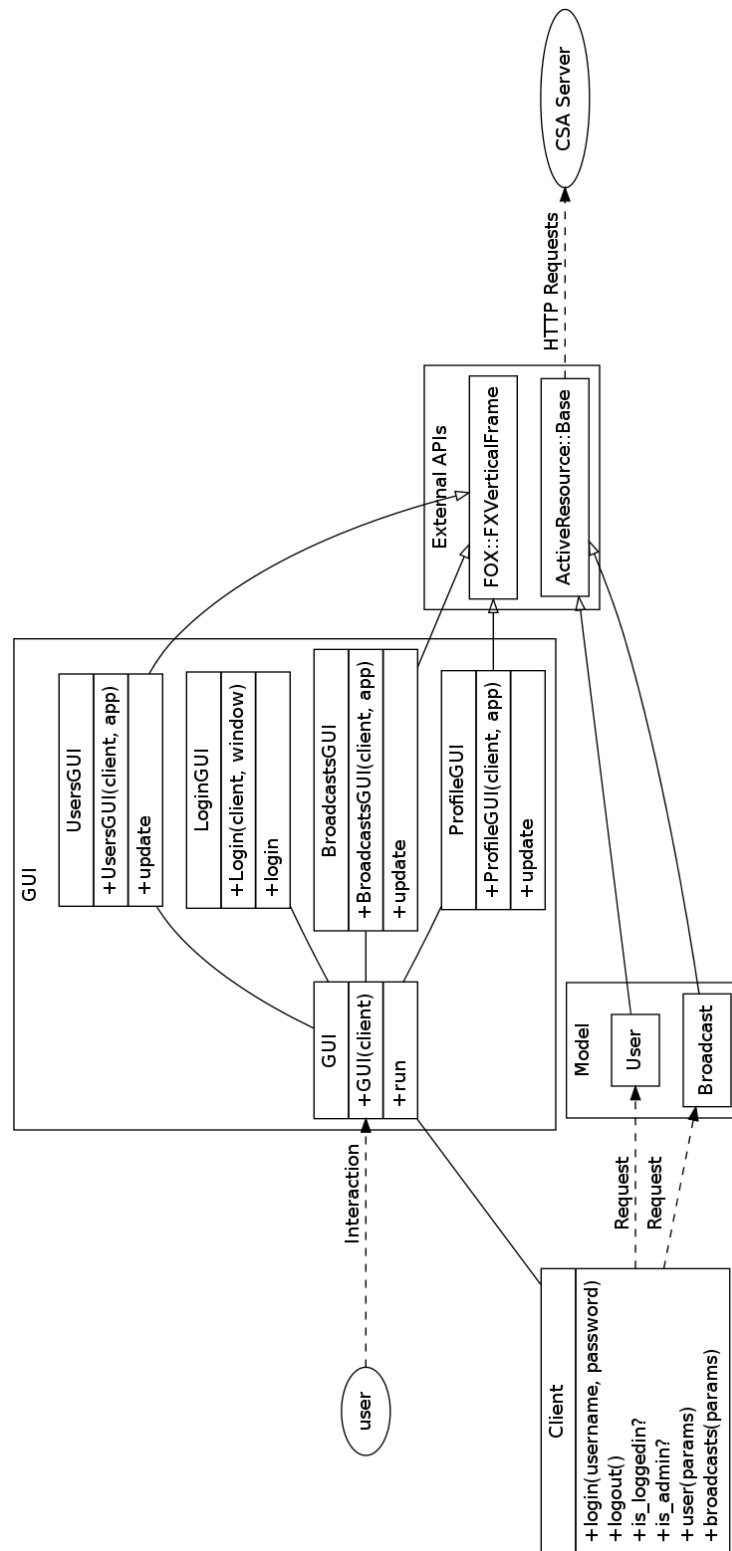


Figure 2: UML Class diagram for the client

## Test Strategy

The first part of the strategy was to fix all the failing tests on the CSA server; the most of these were caused by there not being a logged-in user when performing the actions. It was a simple matter of creating an entry in the `user_details` table in the test fixtures. I then got this entry and added it to a fake session Rails testing provides to assist with this scenario.[1]

With this the majority of the failing tests then passed. The only other broken tests concerned emailing out broadcasts. This was caused by a problem with the code to do this, as mentioned in the Server Architecture section.

With this complete I then turned my attention to testing the CSA client. However, when I started to work out what needed testing I found out there was very little. It seemed pointless to test the `ActiveResource` models in the client as they were provided by an external API which should be well tested. The only thing I could test is that the methods for logging-in and checking administrator users to make sure the GUI only enables the correct options.

The final step was to look back at the server. As all the client does is poll the server for information using `ActiveResource` it stands to reason that if the server is correctly tested, the client should work correctly with it. Of course in real life it's never quite that simple, but just running and using the client should be enough to verify that it's working correctly.

I decided to look at `RSpec` first, it seemed like a nice testing framework I have seen used very successfully on diaspora. `RSpec` enables better end-to-end testing acting on the system how a user might. I designed and implemented some tests to check the flow of a user through the web front-end, testing how different modules interacted with each other.



## Evaluation

I found this assignment quite easy in terms of programming. Having spent a year in industry; especially in a role which involved heavily interacting with existing systems and external APIs; this was all old hat to me. I had also taken some time during my year to look and play with Diaspora - a decentralised social network platform written in Ruby on Rails, so I was somewhat familiar with some of the technologies involved too. I very quickly mastered FXRuby; again I have spent quite a bit of time playing with graphics libraries so picking up a new one which has parallels to Java's SWT, was a fairly easy task; and produced a quick mock up of the GUI I would be using for the Client in a single class. I then abstracted out some of the resources into separate classes where appropriate.

The biggest difficulty I had on this assignment was when it came to updating a user's profile from the client. Because of the way ActiveRecord works it sends every single detail about a model object back through the PUT request. This in turn caused the server too result in a HTTP 500 error code as ActiveRecord was trying to update protected attributes such as the `id`, `created_at` and `updated_at` attributes.

Figure 3 shows the details of the request.

```
PUT /users/41.json HTTP/1.1
Content-Type: text/json
Accept: text/json
... HTTP Authentication headers.

{
  "created_at": "2012-11-05T14:27:03Z",
  "email": "cwl39@aber.ac.uk",
  "firstname": "Firstname39",
  "grad_year": 1985,
  "id": 41,
  "jobs": true,
  "phone": "01970 622422",
  "surname": "Test",
  "updated_at": "2012-12-02T12:25:56Z",
  "image": {
    "id": 1,
    "photo_content_type": "image/jpeg",
    "photo_file_name": "P1210623.JPG",
    "photo_file_size": 2087613,
    "photo_updated_at": "2012-08-13T11:14:04Z",
    "user_id": 41
  }
}
```

Figure 3: Example of a PUT request to user with ID of 41

Digging a little deeper into this issue and looking through the log-files Rails produces it turned out the ImageService class was receiving the Hash generated from the JSON file and putting it straight into the User object. This was then causing ActiveRecord to fail validations for changing protected values and returning a HTTP 500 error code to the request.

I tried filtering out the protected attributes using `ActiveRecord.protected_attributes`, however this seemed to only contain certain attributes and not a whole list. I resolved this by iterating the hashes keys and seeing if `ActiveRecord.available_attributes` contained that key. If not it was removed from the hash. This fixed the problem and didn't cause any other problems. I had also tried nullifying the

attributes client side, but this meant ActiveResource sent the request to the wrong place (as it lost the ID to guide it).

I had a good attempt at getting the CSA Server to use SSL. I tried the provided method which appears should work for older versions of Apache/httpd. However due to the way my home system is set-up (I run either Apache2 or nginx) and a general lack of documentation into how to run Rails as a production environment with Apache2 (I didn't even try looking for nginx as I knew it would undoubtedly be more of a challenge to set up) I was unable to get Rails routing through a normal webserver and by extension couldn't get HTTPS working correctly.

I did try running both WEBrick and Mongrel in isolation with the option `config.force_ssl=true` with very little luck. Both complained about the way the request was coming in, which I can only assume is a side effect of not routing through a proper webserver. Either way at this point I decided it would be better to polish the client rather than spend a lot of time trying to get this to work without the potential of a working result.

## References

- [1] G. V. Buren, “Passing sessions and referers in rails functional tests.” <http://garrickvanburen.com/archive/passing-sessions-and-referers-in-rails-functional-tests/>, October 2007.